

INTRO TO PARALLELIZING R CODE

Jacalyn Huband

22 September 2022

AGENDA



Discussion of Parallelizing Code

Examples of Parallel Packages

Running Parallel Jobs on Rivanna

DISCUSSION OF PARALLELIZING CODE

What does it mean to parallelize code?

- Most programs are written to run *serially* – each line is executed, one at a time.
- In parallel code, **independent tasks** are assigned to separate processors so that they can run simultaneously.

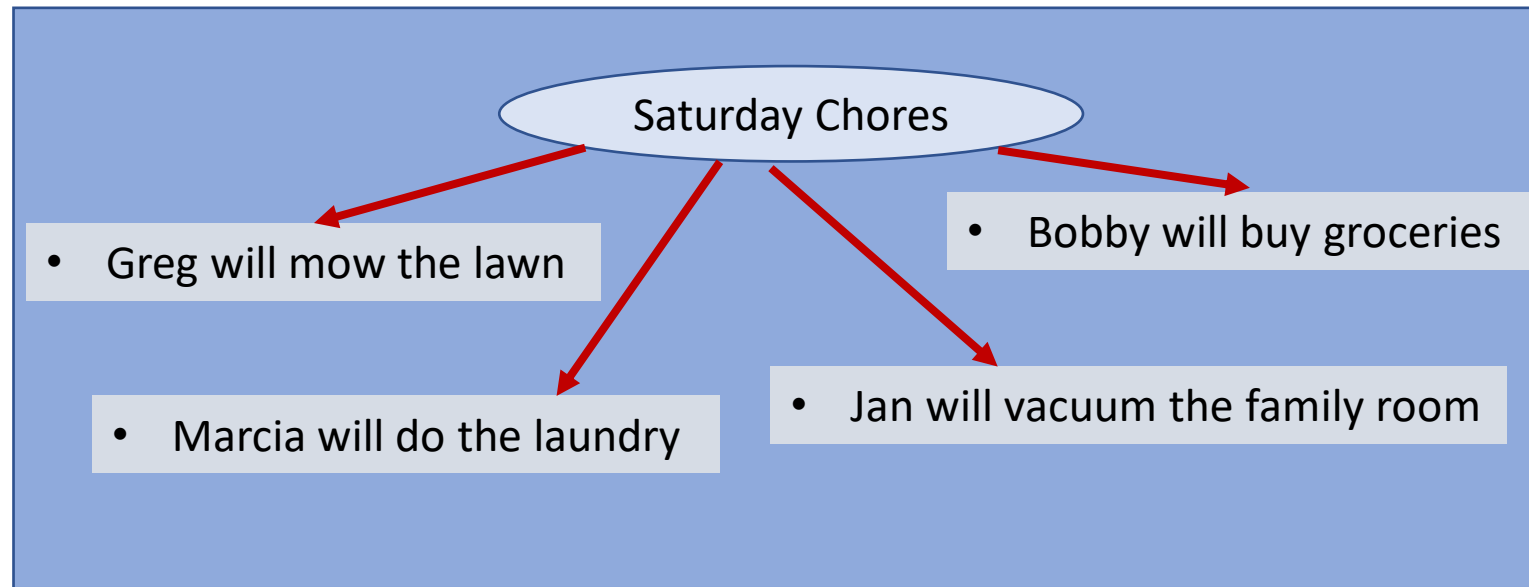
Serial vs. Parallel Concept

• Serial Example



vs.


Parallel Example



Instead of people doing chores, a parallel program assigns tasks to individual processors (cores).

A Note about parallel R

- R uses a **manager/worker** model



- Manager sends tasks to the workers
 - Workers perform the tasks and return the results to the manager
 - The manager combines the results
-
- The parallel packages in R make the communications between the manager and workers transparent to the programmer.

Why parallelize your code?

- Speed

Computers have reached a limit for how fast they can perform a computation.

- Time

When work is being done simultaneously, it can reduce the total amount of time for the code to complete.

Note: If not done properly, parallelizing your code can slow down its performance.

R as a parallel language

- R by itself does not have parallel capabilities.
- There are lots of packages available that can extend R to have parallel capabilities.
 - parallel
 - doParallel
 - multidplyr`
- Today, we will look at a limited number of functions from two packages:
 - *future*
 - *future.apply*
 - *furrr*

Activity

- If you do not have future & furrr installed, go ahead and do that now.

```
install.packages("future")  
install.packages("future.apply")  
  
install.packages("furrr")
```

- We also will be using a timing package tictoc. Install tictoc if you do not have it already:

```
install.packages("tictoc")
```

Activity

- To get a copy of today's scripts, log onto Rivanna and type:

```
cp -r /project/rivanna-training/multicore_R_2022-09-22 ~
```

Or

```
git clone https://github.com/jhuband/Multicore_R.git
```

DETERMINING THE NUMBER OF CORES

The *future* package

What is a core?

- Most computers come with multiple processors or cores.
- On your laptop, the cores allow you to do several tasks at the same time (e.g., typing a report while watching your slack messages and googling the time that your favorite coffee shop closes).
- To take advantage of having multiple cores, you will need to know how many core are available and how to send tasks to the cores

availableCores() vs detectCores()

- **Caution:** Lots of websites tell you to use `parallel::detectCores()` to determine the number of cores.
 - This will report the number of physical cores.
 - On Rivanna, you may not have access to all of the physical cores
- A safer technique is to use `future::availableCores()`
- An even better approach:
`numCores <- max(1, future::availableCores() - 1)`

MULTICORE COMPUTING

The *future* packages

future.lapply package

- A frequently-used functions from the future.apply package:
 - `future_lapply`
 - Allows you to apply the same function to a list of elements
 - Most basic syntax: `results <- future_lapply(x, FUN)`
 - where x = a list of elements, and
 - FUN = a function, either built-in or defined by you
 - results = list of results from the application of FUN to x

multicore Example 1

future_lapply

01_future_lapply.R:

```
library(future.apply)
library(tictoc)
```

#Define the function to be passed to the cores

```
myFunc <- function(value){
  Sys.sleep(1)
  paste0("I received ", value, ". Hello from ",
Sys.getpid())
}
```

#Define the type of parallelization

```
plan("multisession", workers=3)
```

Launch the jobs and collect the results

```
tic()
results <- future_lapply(1:10, myFunc)
toc()
print(unlist(results))
```

Define the task that
will be done in parallel.

Define the type of parallelization.

Launch the parallel
work with
"future_lapply"

multicore Example 1

future_lapply

01_future_lapply.R:

```
library(future.apply)
library(tictoc)
```

#Define the function to be passed to the cores

```
myFunc <- function(i) {
  Sys.sleep(1)
  paste0("I received ", i, " from ")
  Sys.getpid()
}
```

#Define the type
`plan("multisessi`

Launch the jobs and collect the results

```
tic()
results <- future_lapply(1:10, myFunc)
toc()
print(unlist(results))
```

Define the task that

Results:

4.253 sec elapsed

```
[1] "I received 1. Hello from 18341" "I received 2. Hello from 18341"
[3] "I received 3. Hello from 18341" "I received 4. Hello from 18342"
[5] "I received 5. Hello from 18342" "I received 6. Hello from 18342"
[7] "I received 7. Hello from 18342" "I received 8. Hello from 18340"
[9] "I received 9. Hello from 18340" "I received 10. Hello from 18340"
```

Launch the parallel
work with
`"future_lapply"`

An advantage of the future packages

- Because the future package allows you to specify the type of parallelization, you can easily change from serial to parallel, or vice versa.

```
plan("multisession", workers=3)  
plan("sequential")
```

Activity

- Copy 01_future_lapply.R to 01_serial.R
- Change the plan to sequential and run the code.
- What happened to the timing results?
- Is this what you expected?

furr package

- A frequently-used function from the furr package:

- `future_map`

- Allows you to apply the same function to a list of elements
- Most basic syntax: `results <- future_map(x, FUN)`

where x = a list of elements, and

FUN = a function, either built-in or defined by you

results = list of results from the application of FUN to x

multicore Example 2

future_map

02_future_map.R:

```
library(future); library(future.map)
library(tictoc)
```

#Define the function to be passed to the cores

```
myFunc <- function(value){
  Sys.sleep(1)
  paste0("I received ", value, ". Hello from ",
  Sys.getpid())
}
```

#Define the type of parallelization

```
plan("multisession", workers=3)
```

Launch the jobs and collect the results

```
tic()
results <- future_map(1:10, myFunc)
toc()
print(unlist(results))
```

Define the task that
will be done in parallel.

Define the type of parallelization.

Launch the parallel
work with
"future_lapply"

multicore Example 2

future_map

02_future_map.R:

```
library(future); library(future.map)
library(tictoc)
```

#Define the function to be passed to the cores

```
myFunc <- function(i) {
  Sys.sleep(1)
  paste0("I received ", i, " from ")
  Sys.getpid()
}
```

#Define the type of future
`plan("multisession")`

Launch the jobs and collect the results

```
tic()
results <- future_map(1:10, myFunc)
toc()
print(unlist(results))
```

Define the task that
will be done in parallel

Results:

4.615 sec elapsed

```
[1] "I received 1. Hello from 24858" "I received 2. Hello from 24858"
[3] "I received 3. Hello from 24858" "I received 4. Hello from 24859"
[5] "I received 5. Hello from 24859" "I received 6. Hello from 24859"
[7] "I received 7. Hello from 24859" "I received 8. Hello from 24857"
[9] "I received 9. Hello from 24857" "I received 10. Hello from 24857"
```

Launch the parallel
work with
"future_lapply"

DoFuture package

- If you absolutely must have a loop
- The DoFuture package allows you to run a foreach loop in parallel
- You will need an additional step of registering the parallelization

multicore Example 3

DoFuture

03_doFuture.R:

```
library(foreach); library(future); library(doFuture)
library(tictoc)
# Define the function to be passed to the cores
myFunc <- function(value){
  Sys.sleep(1)
  paste0("I received ", value, ". Hello from ", Sys.getpid())
}
# Define the type of parallelization
registerDoFuture()
plan("multisession", workers=3)
# Launch the jobs and collect the results
tic()
results <- foreach(x = 1:10, .combine=c) %dopar%
  myFunc(x)
}
toc()
print(unlist(results))
```

Define the task that
will be done in parallel.

Define the type of parallelization.

Launch the parallel
work with "foreach"
and "%dopar%"

multicore Example 3

DoFuture

03_doFuture.R:

```
library(foreach); library(future); library(doFuture)
library(tictoc)
# Define the function to be passed to the cores
myFunc <- function(value){
  Sys.sleep(1)
  paste0("I received ", value)
}
# Define the type of future
registerDoFuture()
plan("multisession", workers = 4)
# Launch the jobs and get the results
tic()
results <- foreach(x = 1:10, .combine=c) %dopar%
  myFunc(x)
}
toc()
print(unlist(results))
```

Define the task that
will be done in parallel.

Results:

4.328 sec elapsed

```
[1] "I received 1. Hello from 320" "I received 2. Hello from 320"
[3] "I received 3. Hello from 320" "I received 4. Hello from 319"
[5] "I received 5. Hello from 319" "I received 6. Hello from 319"
[7] "I received 7. Hello from 319" "I received 8. Hello from 321"
[9] "I received 9. Hello from 321" "I received 10. Hello from 321"
```

Launch the parallel
work with "foreach"
and "%dopar%"

Activity

I have a program (compute_pi.R) that we can try to parallelize.

```
numDarts <- 1000
circleHits <- 0

for (n in 1:numDarts){
  x <- runif(1); y <- runif(1)
  d <- sqrt(x*x + y*y)
  if (d <= 1.0){
    circleHits <- circleHits + 1
  }
}
# Use formula to estimate pi
pi = 4.0 * circleHits/numDarts
cat("\nThe estimate for pi is",pi,".\n")
```

How would you parallelize this program?

Need more help?

Office Hours via Zoom

Tuesdays: 3 pm - 5 pm

Thursdays: 10 am - noon

To connect to the Zoom sessions,
go to
<https://www.rc.virginia.edu/support/#office-hours> and click on the
“Join us via Zoom” button

Website:

<https://rc.virginia.edu>

Email:

<https://www.rc.virginia.edu/form/support-request/>

Questions?

